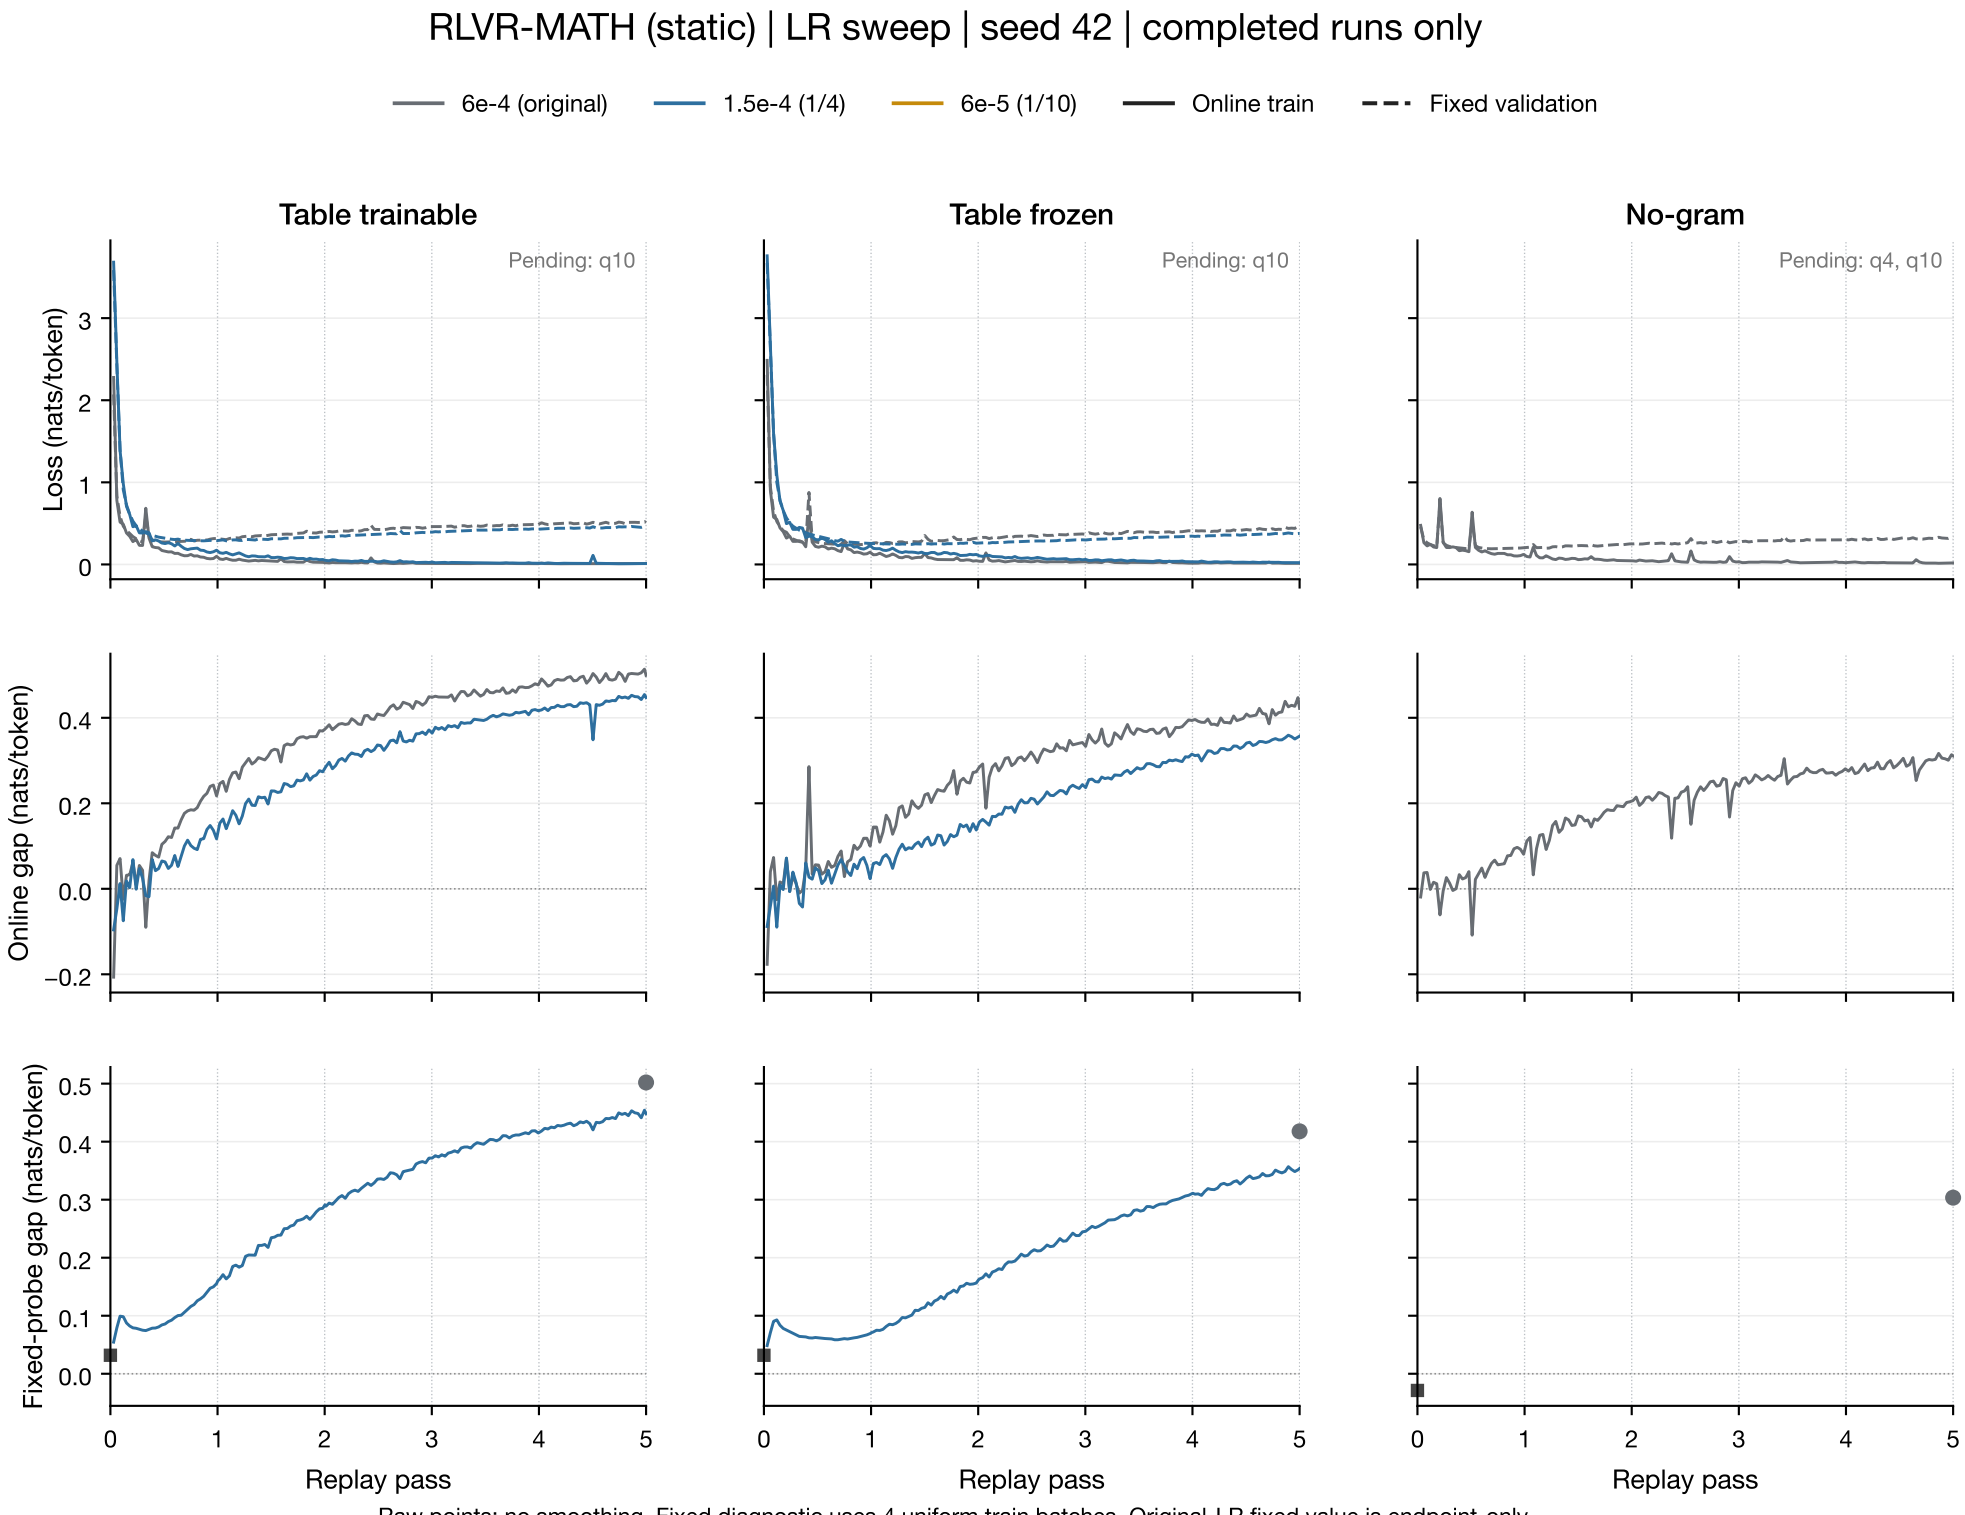
RLVR-MATH (static) | LR sweep | seed 42 | completed runs only
6e-4 (original)
1.5e-4 (1/4)
6e-5 (1/10)
Online train
Fixed validation
Table trainable
Table frozen
No-gram
Pending: q10
Pending: q10
Pending: q4, q10
Loss (nats/token)
Online gap (nats/token)
Fixed-probe gap (nats/token)
Replay pass
Replay pass
Replay pass
Raw points; no smoothing. Fixed diagnostic uses 4 uniform train batches. Original-LR fixed value is endpoint-only.