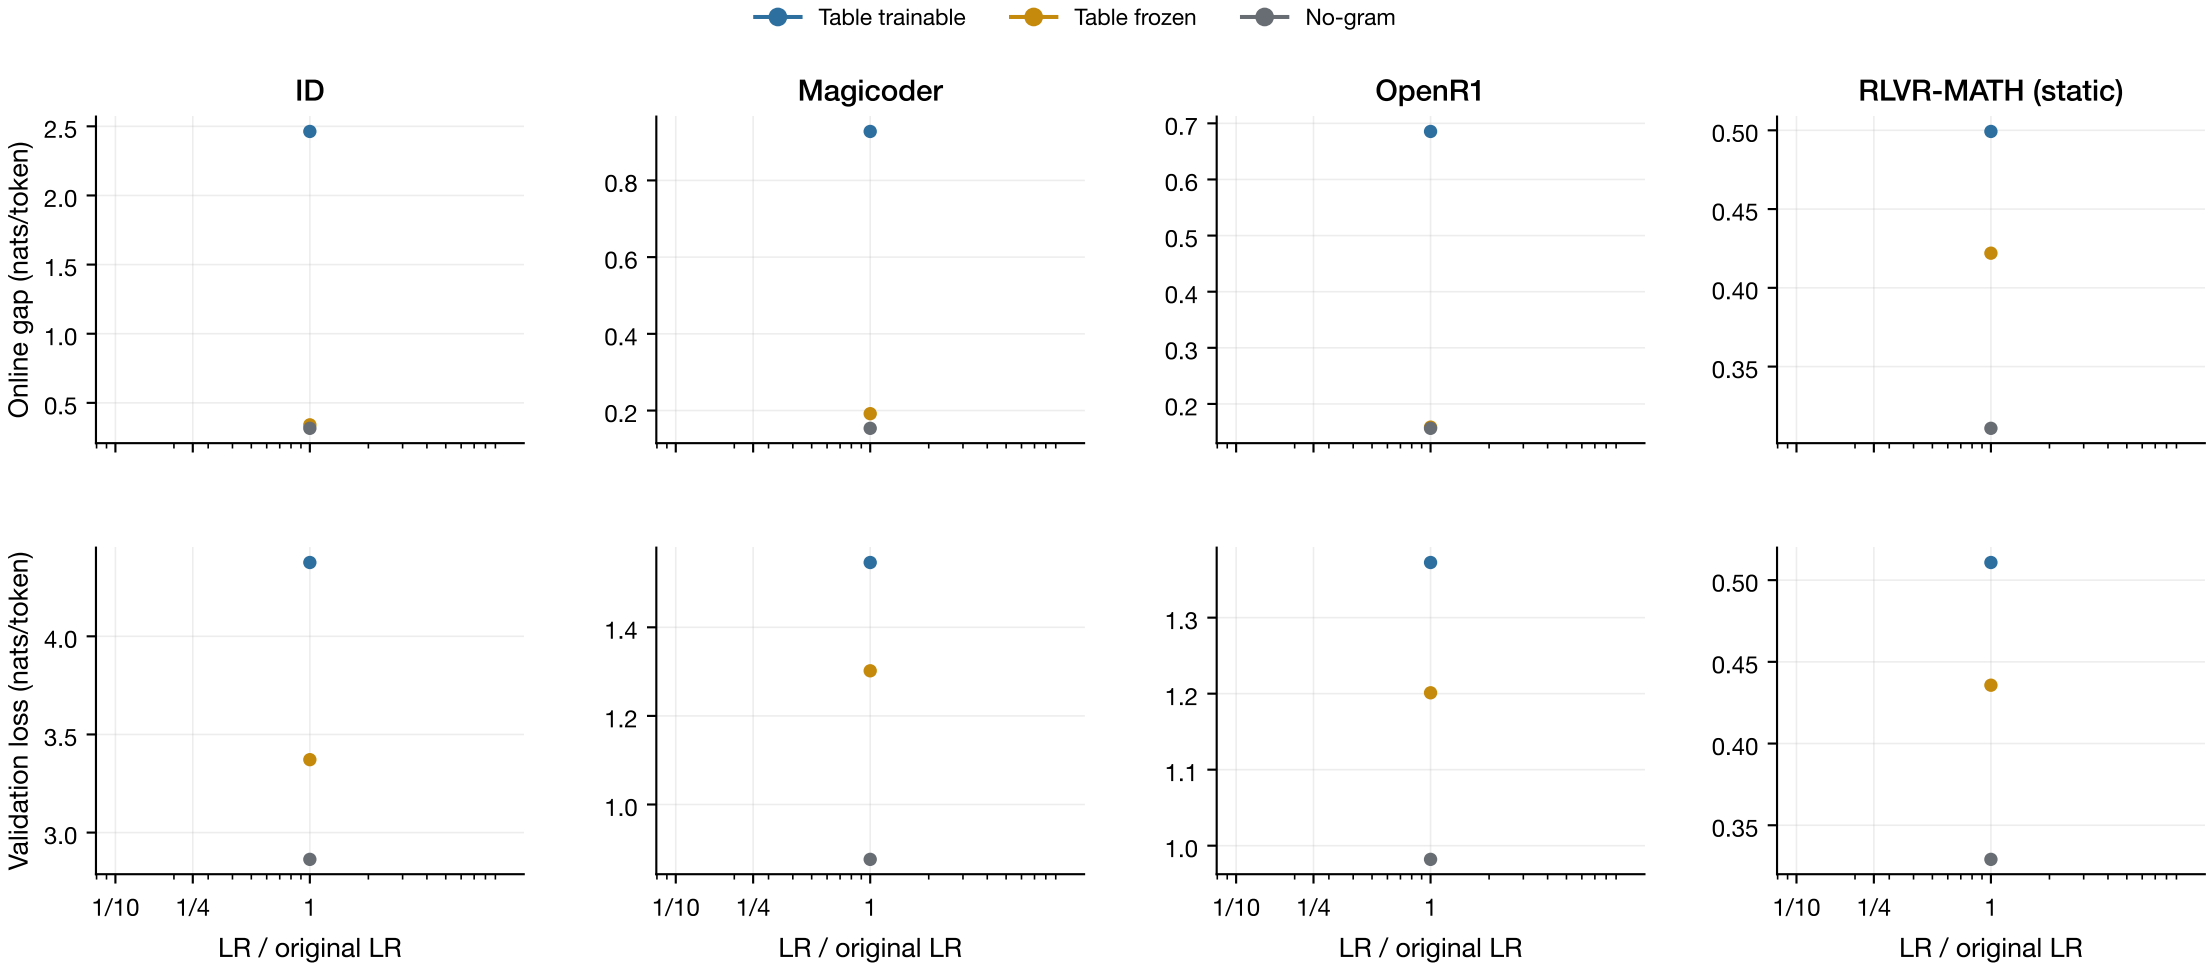
LR sensitivity after five replay passes | seed 42
Table trainable
Table frozen
No-gram
ID
Magicoder
OpenR1
RLVR-MATH (static)
Online gap (nats/token)
Validation loss (nats/token)
1/10
1/4
1
LR / original LR
Same initial checkpoints, data, total tokens and table/backbone LR ratio (128). Lower gap alone does not establish better generalization.