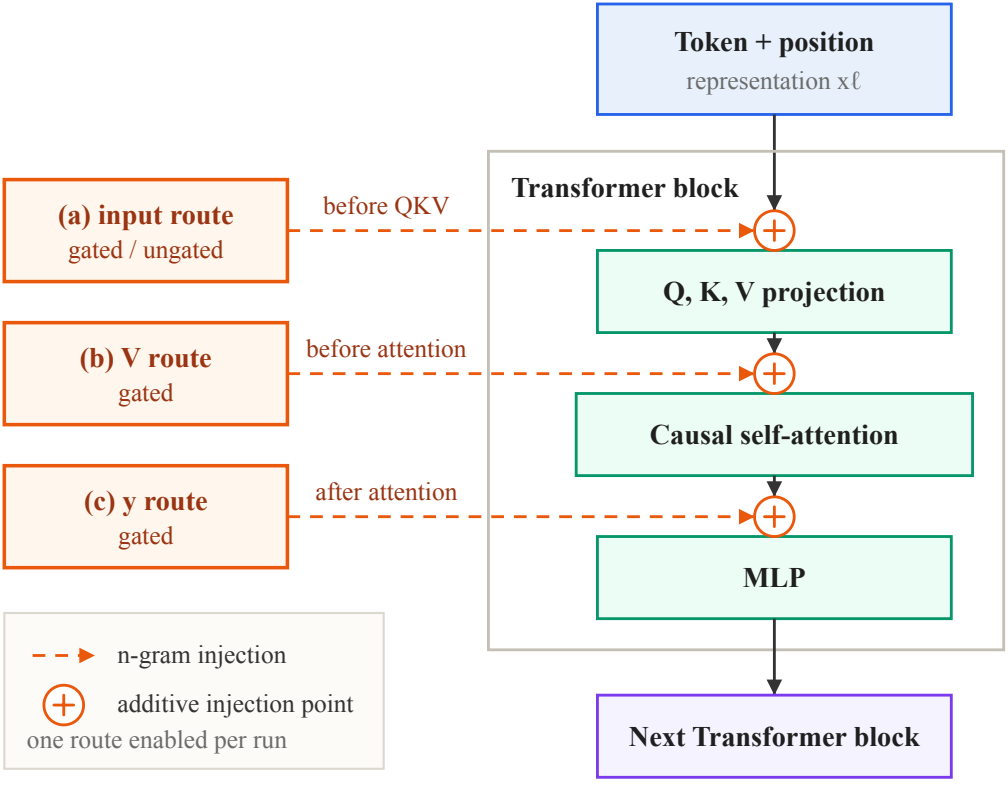
Token + position
representation xℓ
Transformer block
(a) input route
gated / ungated
before QKV
Q, K, V projection
(b) V route
gated
before attention
Causal self-attention
(c) y route
gated
after attention
MLP
n-gram injection
additive injection point
one route enabled per run
Next Transformer block